

ClauseClear AI

Intelligent Legal Contract Analysis Platform

Institution: SRM University AP

Domain: Generative AI · Legal Technology

Team Members:

M Bapu Koushik

J Tharun

V Kishore

Problem Statement

The Core Challenge

- Legal contracts are written in highly complex, technical, and domain-specific language that is inaccessible to the general public
- Non-legal professionals — including individuals and small business owners — frequently sign agreements without fully comprehending the embedded risks or obligations
- This information asymmetry leads to uninformed decisions, legal disputes, and financial harm
- Small enterprises lack the resources to engage legal counsel for routine contract review

Limitations of Existing AI Tools

- Current AI-based solutions produce **generic outputs** that are not tailored to specific legal jurisdictions or contract types
- They lack proper **legal grounding** and do not reference applicable statutes or frameworks
- Susceptibility to **hallucination** — generating plausible but factually incorrect legal information — poses serious risks
- Absence of structured, verifiable output formats makes results unreliable for practical use



There is a critical need for a **reliable, structured, and legally accurate** AI system that bridges the gap between complex legal language and everyday understanding.

Solution Overview

What is ClauseClear AI?

- **ClauseClear AI** is a Generative AI-based web application purpose-built for legal contract simplification and risk analysis
- Users can **upload or paste** contract documents directly into the platform for instant processing
- The system performs multi-dimensional analysis and returns structured, human-readable results within seconds
- Designed to serve individuals, startups, and small businesses who lack legal expertise

What the System Generates

- **Plain-language summaries** — translates dense legal jargon into clear, understandable language accessible to non-lawyers
- **Risk identification** — flags potentially harmful clauses categorized by severity (High, Medium, Low)
- **Structured analysis** — returns results in a consistent JSON format for clarity and integration
- **Indian legal framework alignment** — evaluates contracts against relevant Indian statutes for jurisdiction-specific accuracy



ClauseClear AI transforms an opaque legal document into an **actionable intelligence report** in under 25 seconds.

Objectives

Primary Technical Objectives

- **Develop a multi-stage AI pipeline** encompassing four core phases: Classification (identifying contract type and domain), Analysis (extracting key clauses and obligations), Self-Critique (validating and correcting AI output), and Storage (persisting results for future reference)
- **Implement Retrieval-Augmented Generation (RAG)** to enhance contextual accuracy by injecting relevant historical legal data during analysis
- **Enforce structured JSON outputs** using schema validation to ensure consistent, machine-readable, and verifiable results across all analyses

Performance & System Objectives

- **Maintain session-based conversation memory** to support context-aware follow-up queries, enabling users to ask clarifying questions about their contract analysis
- **Achieve fast processing time** of 15–25 seconds per contract analysis, making the tool practical and efficient for real-world use
- **Minimize hallucination** through a self-critique feedback loop that re-evaluates AI outputs before presenting them to the user
- **Ground all analyses** in applicable Indian legal statutes to ensure jurisdiction-specific accuracy and reliability



These objectives collectively define a **production-grade, legally informed AI system** rather than a generic language model wrapper.

Core Technologies

Backend & AI Stack

- **Python** — Primary backend programming language; chosen for its extensive AI/ML ecosystem, readability, and rapid development capabilities
- **Flask Framework** — Lightweight Python web framework handling API routing, request processing, session management, and response delivery
- **Google Gemini 2.5 Flash Lite** — Core AI engine responsible for natural language understanding, contract analysis, summarization, and structured output generation
- **SQLite Database** — Embedded relational database for persistent storage of analysis results, session data, and conversation history without external dependencies

Frontend & AI Enhancement

- **HTML, CSS, JavaScript** — Standard web technologies used to build a clean, responsive, and intuitive user interface for contract upload and result visualization
- **Vector Embeddings** — Numerical representations of text used to encode contract content and legal reference documents into a comparable mathematical space
- **Cosine Similarity** — Mathematical metric used to identify and retrieve the most contextually relevant historical legal data for RAG-based analysis enhancement
- **JSON Schema Enforcement** — Structural constraint layer ensuring all AI outputs conform to a predefined format, eliminating unparseable or inconsistent responses



The technology stack is deliberately chosen for **simplicity, portability, and zero external vector database dependency**, making the system lightweight and self-contained.

Key Features

Contract Health Assessment

- **Health Score (0–100)** — A quantitative metric indicating the overall quality and fairness of the contract; higher scores denote a lower-risk agreement
- **Letter Grade (A–F)** — An intuitive grading system derived from the health score, providing an at-a-glance risk assessment familiar to all users
- **Risk Classification** — Each identified clause is tagged as **High Risk** (potentially harmful obligations), **Medium Risk** (ambiguous or unfavorable terms), or **Low Risk** (standard, acceptable provisions)
- **Obligations Extraction** — Automatically identifies and lists all binding duties and responsibilities imposed on each party

Legal Playbooks & Accuracy

- **Self-Critique Loop** — After initial analysis, the AI re-evaluates its own output against the source contract to identify and correct errors before delivery
- **Indian Contract Act (1872)** — Governs offer, acceptance, consideration, and enforceability standards applied during clause evaluation
- **DPDP Act (2023)** — Digital Personal Data Protection Act provisions applied to identify data privacy obligations and non-compliant clauses
- **IT Act (2000)** — Information Technology Act provisions applied to assess electronic contract validity and cybersecurity-related terms



The integration of three Indian legal frameworks ensures **jurisdiction-specific, legally grounded analysis** rather than generic AI commentary.

AI Pipeline — How It Works

Input & Context Preparation (Steps 1–4)

- **Step 1 — Contract Ingestion:** User uploads a contract file (PDF, DOCX, TXT, or image) or pastes raw text directly; the system extracts and normalizes the content into plain text
- **Step 2 — System Prompt Engineering:** A carefully designed system prompt defines the AI's role as a legal analyst, sets the evaluation framework, and constrains the scope of response
- **Step 3 — JSON Schema Enforcement:** A predefined output schema is injected into the AI request, ensuring the model returns data in a structured, consistent, and parseable format
- **Step 4 — RAG Context Injection:** Vector embeddings of the contract are compared against stored legal references using cosine similarity; the most relevant legal context is retrieved and appended to the AI prompt

Analysis & Output Delivery (Steps 5–7)

- **Step 5 — Core AI Analysis:** Google Gemini 2.5 Flash Lite processes the enriched prompt — performing clause identification, risk classification, obligation extraction, and plain-language summarization simultaneously
- **Step 6 — Self-Critique Validation:** The generated analysis is passed through a second AI call that critiques the output for accuracy, completeness, and legal consistency; corrections are applied automatically
- **Step 7 — Result Delivery:** The validated, structured JSON result is rendered on the frontend displaying the health score, grade, risk-tagged clauses, obligations, and plain-language summary to the user



The self-critique loop in Step 6 is the key differentiator that **significantly reduces hallucination** compared to single-pass AI systems.

System Architecture

Layered Architecture Design

- **Frontend Layer** — Built with HTML, CSS, and JavaScript; provides the user interface for contract upload, text input, and structured result visualization including the health score dashboard and risk-tagged clause viewer
- **Backend Layer** — Flask application server handles all incoming API requests, manages file parsing (PDF/DOCX/TXT), orchestrates the AI pipeline, and delivers structured JSON responses to the frontend
- **AI Pipeline Layer** — The intelligence core; responsible for context engineering, prompt construction, RAG retrieval, Gemini API invocation, self-critique execution, and output validation

Data & Design Principles

- **Database Layer** — SQLite stores all analysis results, session identifiers, conversation history, and legal reference embeddings; enables persistent memory and retrieval without external database infrastructure
- **Modular Design** — Each layer operates independently with well-defined interfaces, allowing components to be updated, tested, or replaced without affecting the entire system
- **Scalable Architecture** — The stateless Flask API design allows horizontal scaling; the SQLite layer can be replaced with PostgreSQL for production-scale deployments
- **No External Vector DB** — RAG is implemented natively using in-memory vector computations, eliminating dependency on external services like Pinecone or Chroma



The architecture prioritizes **simplicity, self-containment, and deployability** — the entire system can run on a single server without cloud dependencies.

Results & Performance

Performance Benchmarks

- **Analysis Time:** Approximately **12–22 seconds** per full contract analysis, inclusive of RAG retrieval, AI processing, self-critique loop execution, and JSON rendering — well within the target range of 15–25 seconds
- **JSON Output Accuracy:** Approximately **98%** structurally valid JSON outputs produced, meaning near-complete elimination of unparseable or malformed responses in testing
- **Self-Critique Effectiveness:** The validation loop successfully identifies and corrects factual inconsistencies, missing fields, and hallucinated clauses before output delivery

Compatibility & Memory

- **Supported File Formats:** Accepts **PDF** (most common contract format), **DOCX** (Microsoft Word documents), **TXT** (plain text contracts), and **images** (scanned contracts via OCR processing)
- **Contract Length Handling:** Processes contracts up to **15,000 characters** — sufficient for most standard commercial and employment agreements — with graceful truncation for oversized inputs
- **Session Memory:** Maintains full conversation history per session, enabling users to ask context-aware follow-up questions such as "explain clause 4 in simpler terms" without re-uploading the document
- **Concurrent Sessions:** Flask's session management supports multiple simultaneous user sessions with isolated memory contexts, ensuring data privacy between users



The system achieves a strong balance between **speed, accuracy, and legal reliability** — validating the core architectural decisions made during design.

Advantages & Future Scope

Key Advantages

- **Legally Grounded Analysis:** Unlike generic AI tools, ClauseClear AI references the Indian Contract Act (1872), DPDP Act (2023), and IT Act (2000), providing jurisdiction-specific and statute-backed evaluations
- **Reduced Hallucination:** The self-critique mechanism validates AI output against the source document, significantly reducing the risk of fabricated legal interpretations
- **Improved Contextual Accuracy via RAG:** Retrieval-Augmented Generation enriches prompts with relevant historical legal context, resulting in more precise and relevant clause assessments
- **Zero External Vector Database Dependency:** RAG is implemented natively, keeping the system self-contained, cost-efficient, and easy to deploy
- **Simple, Efficient Architecture:** A clean four-layer design ensures maintainability, testability, and the ability to scale or extend the system with minimal refactoring

Future Scope

- **International Legal System Support:** Extend the legal playbook library to include GDPR (EU), UCC (USA), and other international frameworks, enabling ClauseClear AI to serve a global user base
- **AI-Based Contract Negotiation Assistant:** Develop an intelligent negotiation module that not only identifies risky clauses but suggests legally sound alternative language for renegotiation
- **Mobile Application Development:** Build native iOS and Android applications to make on-the-go contract analysis accessible to field professionals and remote users
- **Real-Time Contract Editing and Highlighting:** Introduce an inline document editor that highlights risky clauses directly within the contract text, enabling users to make edits in real time with AI-guided suggestions



ClauseClear AI represents a **foundational step toward democratizing legal literacy** — making contract intelligence accessible to everyone, not just legal professionals.